

A Convolutional Neural Network Solution For Facial Expression Classification

Alexandru Cohal
March, 2018

Prof. Dr.-Ing. Andreas König

Overview

1. Motivation
2. Introduction
3. Objective
4. Data Acquisition
5. Architecture
6. Performance
7. Hyperparameters
8. Results
9. Conclusion

Motivation

- The world's older population is growing (in present, approximately 8.5% of global population is aged *65 and over*)
- High number of persons with diminished abilities (*1 out of 5* persons has hearing issues, *1 out of 4* persons has vision issues, *1 out of 4* persons has pains in wrists / hands)
- In the USA in 2013, 17.6% of people with a disability were employed
- The general current trend ➡ **make technology available to more users with different needs and goals** ➡ simplify people's lives and sustain the economy


<http://careforyou.us/making-technology-available-to-everyone/>

<https://www.dosomething.org/facts/11-facts-about-physical-disability>

Introduction

- A **Brain – Computer Interface (BCI)** ➡ a solution for *Human – Computer Interface* ➡ uses signals captured from the brain for controlling an external activity
 - **Electroencephalography (EEG)** ➡ a method for monitoring the electrical neural activity of the brain
 - **Electromyography (EMG)** ➡ monitoring the electrical activity of the muscle tissue ➡ facial expressions can be detected


- Blink
- Left / Right wink
- Raise brow
- Smile



<https://jeffsearle.blogspot.de/2015/04/muscles-of-head-and-neck.html>

Introduction

- **Emotiv EPOC** ➔ a compact, wireless EEG device, easy to setup and use
- Has 16 sensors (electrodes): 14 for data (μV) and 2 for reference




- Data is acquired with 128 Hz frequency
- A saline solution has to be used for increasing conductivity

Investigating the Emotiv EPOC for cognitive control in limited training time – M. Lang, T. Mitrovic, *University of Canterbury* (2012)
<https://www.emotiv.com/epoc/>

Introduction

- **Deep Learning** ➡ a **Machine Learning** family of methods for learning data representations (discover useful features for classification)
- It is composed of multiple processing layers ➡ used to learn representations of data with multiple levels of abstraction
- These methods have obtained very good results ➡ speech recognition, visual object recognition, object detection, etc.

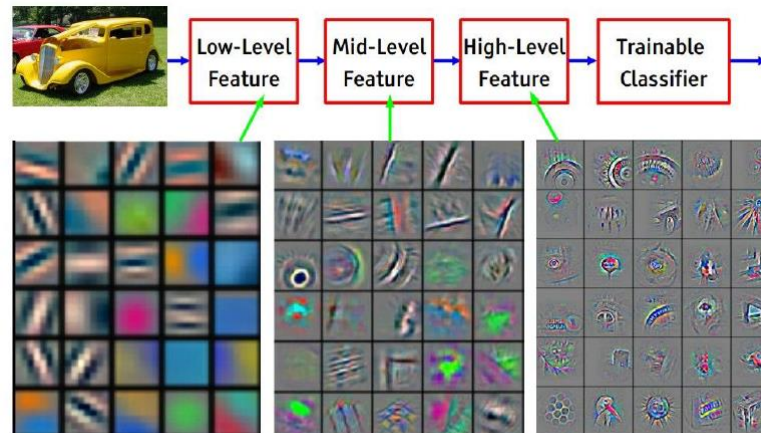


Deep Learning - Y. LeCun et al., <http://www.bioinfo.org.cn/~casp/temp/DeepLearning.pdf>

ImageNet Classification with Deep Convolutional Neural Networks – A. Krizhevsky et al., Advances in Neural Information Processing Systems 25 (2012)

Introduction

- **Classic classification** → manual engineering of features → fed into a machine learning algorithm → high understanding of the domain is required
- **Deep Learning** → does not require manual engineering of features → it learns during training multiple “filters” with increasing complexity as the layers get deeper → uses them in a final classifier




Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

Deep Learning - Y. LeCun et al., <http://www.bioinfo.org.cn/~casp/temp/DeepLearning.pdf>
<https://burakhimmetoglu.com/2017/08/22/time-series-classification-with-tensorflow/>

Introduction

- **Convolutional Neural Networks (CNN)** → is a class of Deep Learning
- It is similar to a classic Artificial Neural Network → only that it exploits spatially-local correlation
- A **convolutional layer** → convolves the data with multiple small kernels (rolls the filter over the data) → these kernels are learned, not predefined
- By stacking multiple convolutional layers → more and more complex features can be learned




<http://deeplearning.net/tutorial/lenet.html>

“Computer Vision” course – F. Lindseth, NTNU, 2017

Introduction

- **Convolutional Neural Networks (CNN)** → different layers:
 - **Input Layer** – holds the raw input data
 - **Convolution Layer** – computes the convolution (output of neurons) between the data and the kernel → multiple kernels are used → the output is a volume, not only a single layer
 - **ReLU Layer** (usually included in the Convolution Layer) – applies the ReLU activation function → introduces non-linearity
 - **Pooling Layer** – downsamples along the spatial dimensions → a smaller volume results
 - **Fully connected layer** – computes the score for each class



“Computer Vision” course – F. Lindseth, NTNU, 2017
<http://cs231n.github.io/convolutional-networks/>
<https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

Introduction

- **Tensorflow** → an open-source machine learning framework
- Uses data flow graphs → the user's written code describes a computation graph → the graph fully describes the desired computation
- It is a portable framework → can run on multiple platforms: CPUs, GPUs, mobile, embedded
- Has interfaces for *C++*, *Java* and *Python*



<https://www.tensorflow.org/>

<https://opensource.com/article/17/11/intro-tensorflow>

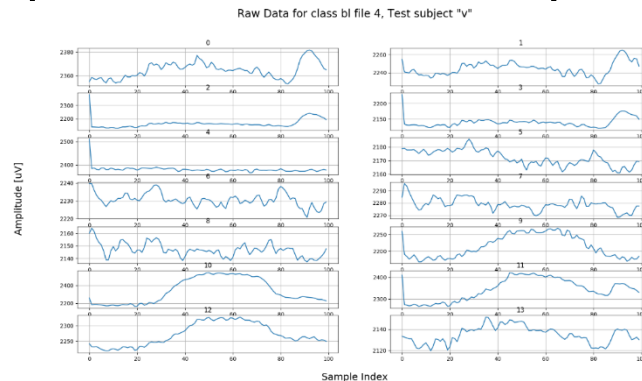
Objective

- Using the *14* channels data provided by the *Emotiv EPOC* headset, classify using a Convolutional Neural Network (CNN) the *Neutral* state and 5 facial expressions:
 - Left wink / blink (BL)
 - Right wink / blink (BR)
 - Strong blink (BB)
 - Open mouth (OM)
 - Full mouth (FM)



Data Acquisition

- For each class:
 - 20 recording of 100 samples each (0.78 sec) → the length was observed to be enough for capturing the whole behavior of the considered facial expressions
 - Each recording contains data from all the 14 channels (sensors)
 - Only from one person → different persons have different behavior of the same facial expressions
 - Only from one session of recording → data is influenced by the position of the headset, quantity of saline solution used



Example of data acquired while performing a “Left Blink” expression

Architecture

- The data acquired from each sensor of the headset is *1D* → **Time-series CNN** is used → the kernels (filters) used are *1D* (not *2D* as in the case of images)
- At each sampling time → Each sensor of the *Emotiv EPOC* headset returns a single values
- The headset has *14 sensors* → at each sample the headset returns an array of *14 values*
- The training is not performed on data from individual time instances (it is irrelevant, can be affected by noise) → windows are used → the input of the CNN will be a *2D* array

Architecture

- The (general) architecture of the CNN is:

<i>Layer</i>	<i>Data Size</i>
Input layer	32×14
1D Convolution layer	32×28
1D Pooling layer	16×28
1D Convolution layer	16×56
1D Pooling layer	8×56
1D Convolution layer	8×112
1D Pooling layer	4×112
1D Convolution layer	4×224
1D Pooling layer	2×224
Fully connected layer	1×6

Convolution layer: each filter has size 2 and moves with stride 1

Pooling layer: the pooling size is 2 and the stride is 2

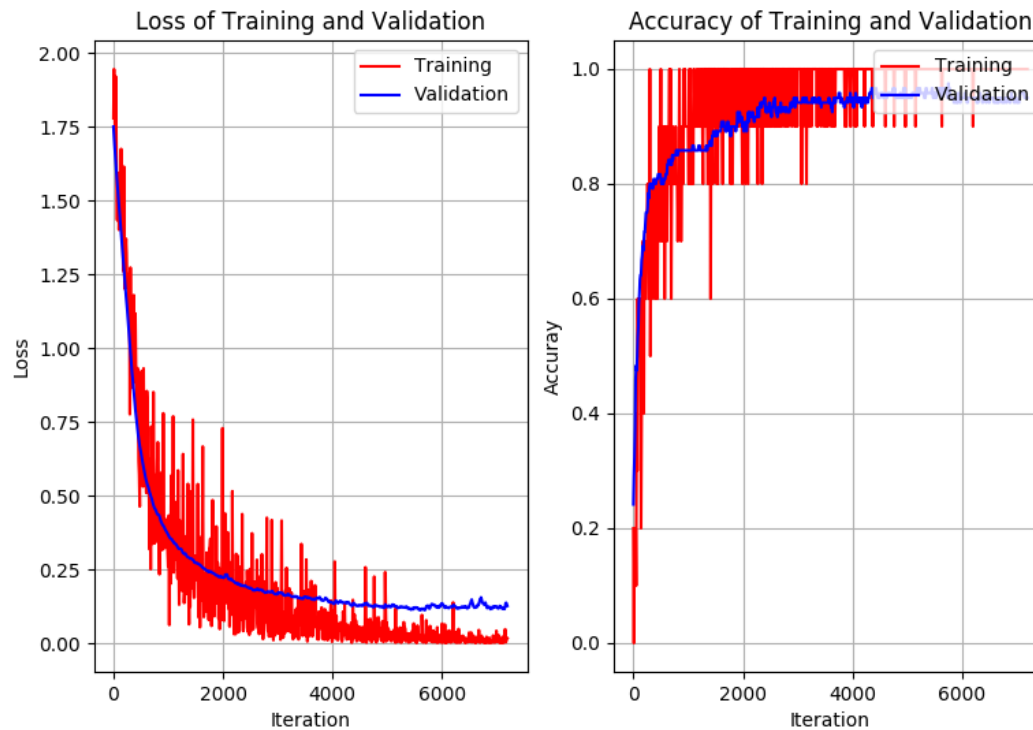
<https://burakhimmetoglu.com/2017/08/22/time-series-classification-with-tensorflow/>

Performance

- The performance of the CNN can be evaluated through:
 - **Loss** → used to optimize the parameters of the network → sum of errors made for each example → better to be as close as possible to 0
 - Loss of training dataset
 - Loss of verification dataset
 - **Accuracy** → used to evaluate how good are the predictions after the network's parameters are optimized
 - Accuracy of training dataset
 - Accuracy of verification dataset
 - Accuracy of testing dataset
 - **Training time** → if the training has to performed multiple times, it is better to be as small as possible

Performance

Testing Accuracy: 91.66 %; Training time: 58.24 s



Hyperparameters

- What can be changed in the CNN in order to obtain better results:
 - ***Training / Validation / Testing datasets ratios***
 - *Training* → the dataset used to adjust the parameters (weights) of the network
 - *Validation* → used to estimate how well the network has trained while training → used for tuning the network hyperparameters (e.g. number of hidden neurons), for early stopping, for observing overfitting
 - *Testing* → the dataset used after the training is done, to prove the predictive capacity of the network
 - The tests performed used the same ratios:
 - Training: 60%, Validation: 20%, Testing: 20%


Hyperparameters

- What can be changed in the CNN in order to obtain better results:
 - *Window size and Overlapping size of windows*
 - Larger windows → too much information, too less samples for training
 - Smaller windows → too few relevant information (learn noise), too many samples for training, longer training time, not enough time for predicting in real-time
 - Larger overlapping → too many irrelevant samples, longer training time, overfitting
 - Lower overlapping → too few samples, loose the behavior from the border
- Window sizes used: 16 and 32 samples
- Overlapping size used: 0, 8, 16 and 24 samples

Hyperparameters

- What can be changed in the CNN in order to obtain better results:
 - *Window size and Overlapping size of windows*


Testing Accuracy: 98.0 %; Training time: 31.21 s



Too small window (16), long training time (400 epochs are not enough)

Hyperparameters

- What can be changed in the CNN in order to obtain better results:
 - *Artificial input data generation*
 - For each window from the original dataset, multiple windows can be obtained by adding noise
 - Too many artificial data used \longrightarrow overfitting, learn noise
 - Less artificial data used \longrightarrow too less data for training
 - Normal distributed noise with mean 20 and standard deviation 10 was used in order to increase the available data 2, 3 and 4 times



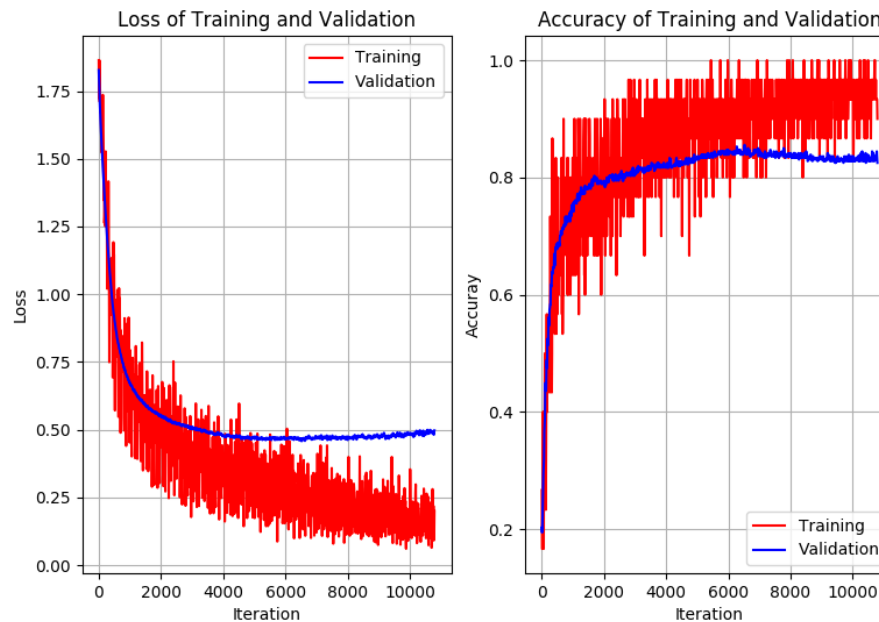
https://en.wikipedia.org/wiki/Normal_distribution

Hyperparameters

- What can be changed in the CNN in order to obtain better results:

- *Artificial input data generation*

Testing Accuracy: 83.05 %; Training time: 121.69 s



Too many artificial data used (3 times more than the original data), Overfitting


Hyperparameters

- What can be changed in the CNN in order to obtain better results:
 - *Batch size*
 - *Batch* → the number of training samples used for one update of the network's parameters
 - Too larger batches → few training batches, computational expensive
 - Too small batches → many training batches, slow, noisy variation of the parameters
 - Batch sizes used: between 10 and 100 samples

Hyperparameters

- What can be changed in the CNN in order to obtain better results:

- *Batch size*



Too small batches (10 samples), Noisy behavior, Overfitting

Hyperparameters


- What can be changed in the CNN in order to obtain better results:
 - *Number of epochs*
 - Epoch → the number of times the networks sees the entire training dataset → different from *Iteration* → the number of times the networks sees a batch
 - Too many epochs → overfitting
 - Too few epochs → underfitting
 - Number of epochs used: 200, 300, 400, 500 and 600

Hyperparameters

- What can be changed in the CNN in order to obtain better results:


- *Number of epochs*

Testing Accuracy: 91.66 %; Training time: 18.96 s



Underfitting (200 epochs)

Testing Accuracy: 84.16 %; Training time: 201.44 s



Overfitting (400 epochs)

Hyperparameters

- What can be changed in the CNN in order to obtain better results:
 - *Learning rate*
 - *Learning rate* → the size of the “step” done in the direction of the negative gradient
 - Too large learning rate → oscillations
 - Too small learning rate → slow convergence, stuck in local minima
 - Learning rate used: *0.0001*


Hyperparameters

- What can be changed in the CNN in order to obtain better results:
 - *Network structure*
 - Change of number of layers
 - Change of layers' sizes
 - Too less layers → loose the advantage of Deep Learning (learn representations of data with multiple levels of abstraction)
 - Too many layers → overfitting, vanishing gradient
 - Number of layers used: 1, 2, 3 and 4
 - Layers' sizes were changed when the window size was changed

Hyperparameters

- What can be changed in the CNN in order to obtain better results:

- *Network structure*




Too many layers (4), Overfitting

Results

- **The best result obtained**


- 4 Layers, Window size: 32 samples, Overlapping: 16 samples, No artificial data, Batch size: 30 samples, Epochs: 300
- Testing accuracy: **95%**

Testing Accuracy: 95.0 %; Training time: 41.86 s

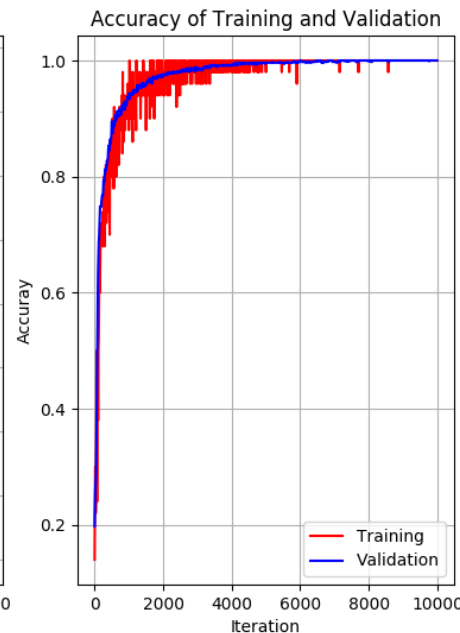
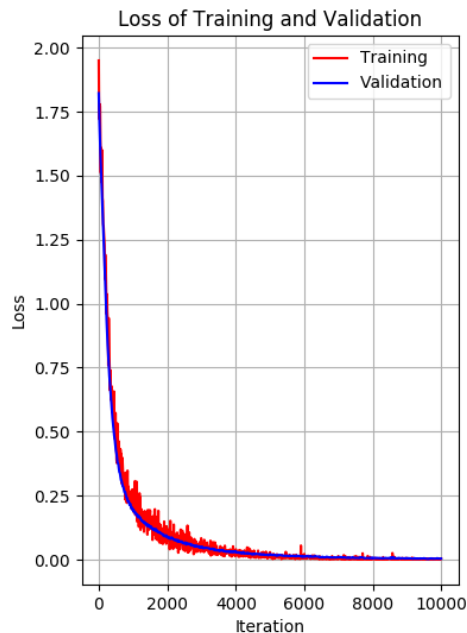


Results

- **Better than the best result obtained**

- 3 Layers, Window size: 32 samples, Overlapping: 30 samples, No artificial data, Batch size: 50 samples, Epochs: 200
- Testing accuracy: **100%** 

Testing Accuracy: 100.0 %; Training time: 156.32 s



*The evolutions of
loss and accuracy
are smoother*

Not replicable



*“Some things in
life only happen
once, the memories
of them lasting
forever”*

J.M.Darhower

Conclusions

- **Time-series Convolutional Neural Network (CNN)** → can be used for classification of facial expressions based on the signals acquired from the *Emotiv EPOC* headset
- The accuracy obtained was very good → 95% (even 100% once) for the testing dataset
- For this case → not a very large dataset was required (20 recordings of 0.78 seconds each for every class) and the training was not long (42 seconds)
- These results → only for one person
- The possibility of generalization could be studied → a large dataset will be required (data from multiple persons and for different positions of the headset, quantities of saline solution used, etc.)

